

Nobody Told the Census It Was Doing Psychology

A Wait But Why-style take on two kinds of measurement

Sep 16, 2026 · @Someone

In Brazil's PNAD Contínua, one person in the house can answer for everyone in it. The person responsible for the household, or whichever adult resident happens to be home, reports the age, schooling, working hours and employment status of people who are out at work and have no idea this is happening.

Your brother-in-law, who has never asked you a direct question about your life, is describing your labor market situation to the federal government.

Psychologists have a name for this. They call it informant rating, and they have been arguing about its reliability since the 1980s. Which is strange, because the census is supposed to be the instrument that deals in hard external facts, and personality testing is supposed to be the soft one that asks how you feel about parties.

The differences are real

The two instruments look like opposites, and on the surface they are.

A census is addressed to a place: "the resident of" a housing unit, whoever that turns out to be this decade. A personality test is addressed to one person and follows them around for life.

A census asks what your walls are made of, how many rooms you have, whether water arrives through a pipe, what you earned before tax and after. The Big Five asks you to agree or disagree with sentences about yourself. One set of answers you could in principle photograph, and the other exists only as a report about an interior.

They also end up in different places. A census answer helps decide where a clinic gets built and how much money a Brazilian state receives through the FPE transfer. A Big Five profile produces a coaching conversation, a clinical note, or a seat on a team that somebody is assembling.

None of that is in dispute. It stops being the interesting part once you look at how the two instruments are put together.

Neither poverty nor conscientiousness exists in nature

There is no fact of the matter about who is poor. There is only a definition somebody wrote down. The Official Poverty Measure compares pre-tax money income against a fixed threshold. The Supplemental Poverty Measure counts SNAP, subtracts taxes and medical and work expenses, and adjusts for what housing costs in the place you actually live. That produces two different poverty rates for the same country in the same year, and neither one was discovered. Both were constructed, for stated reasons.

Employment status works the same way. You count as employed if you worked for pay for at least an hour in the reference week, and that answer comes out of a chain of yes/no questions routing you down one branch of the form rather than another. Whether you are employed is produced by skip logic.

Conscientiousness is built out of comparable decisions. Somebody chose the statements, chose the scale, decided which responses cluster together and what to call the cluster. The score is real in the sense that it predicts things, and not real in the sense of having been sitting there waiting to be found.

Both instruments do the same job: take something abstract and contested, write down explicit rules for turning it into a number, then sort millions of people by those rules. Which is why "is the Big Five valid?" and "is the OPM the right poverty line?" turn out to be closer to one question than they sound.

Both also force answers into closed boxes so they can be added up. The difficulty-seeing item runs from a lot to a little to none at all; the Big Five runs from strongly agree to strongly disagree. Neither instrument offers a box for "it depends," because a million of those aggregate into nothing.

The difference that survives

Both methods share a weakness. People don't respond, and the people who don't respond aren't a random sample of people. The Current Population Survey's ASEC supplement fell to a 61.3% response rate in 2026, and every point of that is an opportunity for the picture to drift toward whoever was reachable.

What matters is that the census knows this about itself. It applies survey weights and control totals precisely because it assumes the sample is wrong in a particular direction, and it publishes the response rate so you can judge how hard that correction is working.

A personality test taken online mostly does none of this. The norms your score is compared against came from whoever showed up, which has historically meant a great many psychology undergraduates and a great many people who enjoy taking personality tests. Nobody weights an Openness score for the fact that incurious people were less likely to click the link.

The difference worth defending is institutional rather than philosophical. One instrument has to justify its numbers in public, which forces it to write its assumptions somewhere a critic can reach them. The other usually doesn't have to, so it usually doesn't.

Three notes on the source framing

Not part of the post. Flagging these before the material reaches students.

The objective/subjective contrast is weaker than it looks. Race and color are self-identified on both the US and Brazilian instruments. Income is self-reported and systematically misremembered. A large share of household answers come from a proxy rather than the person being described. The contrast the evidence supports is documented versus undocumented, which is what the closing section argues instead.

Proxy reporting is the strongest item on the list. It works best in first position because students don't see it coming, and it does the work of unsettling the objective/subjective split before the split gets stated.

Skip logic defining employment is the most useful classroom example. It's concrete, checkable against the actual questionnaire, and it makes the act of construction visible in a way that arguing about the Big Five rarely does.